

The Limitless Builder

Where Imagination Meets
Hermes-Led Vibe Coding

Joe Chow • @matlabulous • Numerai CoE

HOW


Hermes AI Agent + Vibe Coding

- **Hermes** is an AI orchestration agent developed by Nous Research
- Persistent memory + self-improving
- Open source. Connect to any models.
- I talk to it via Terminal / Telegram
- Agent plans, delegates, builds, and reports
- **Vibe Coding** — describe what you want in plain language, AI handles the coding
- Your role shifts from “developer” to “director” — you describe, review, iterate




HERMES-AGENT

Game



Website



councilofelders.github.io

Skill: “plan-event”

Decentralized AI Day - Warsaw - Research - 2026 Q4

Warsaw (warsaw-pl) — Run A: Venue research

Target attendance 50 Soft budget USD 5000

Venue matrix — eligible (Gold) and shortlist-ready (Silver) venues; leads appear only in the always-visible non-selectable list below. Event-duration costs are shown for the select duration.

Venue	Tier	Score	Cost USD	Capacity	Screen	Wi-Fi	Catering	Recording	Event window	Flags	Booking li
ADN Centrum Konferencyjne	Silver	0.5867	?	75	✓ confirmed	✓ confirmed	✓ in_house_available	? unknown	? unknown	•cost_unknown •availability_unconfirmed •final_quote_required •duration_unconfirmed •map_location_unknown •stale_evidence	book / contact
WE EVENT — Sala Eventowa Warsaw Trade Tower	Silver	0.5617	?	80	? unknown	✓ confirmed	✓ in_house_available	? unknown	? unknown	•cost_unknown •screen_unknown •availability_unconfirmed •final_quote_required •duration_unconfirmed •map_location_unknown •stale_evidence	book / contact
Messalka Events	Silver	0.5367	?	80	? unknown	? unknown	✓ in_house_available	? unknown	? unknown	•cost_unknown •screen_unknown •wifi_unknown •final_quote_required •duration_unconfirmed •map_location_unknown •stale_evidence	book / contact

Skill: “create-content”

The prompt – copy from here
(1/3)

09:26 PM

copy

Create ONE image: a minimal world map infographic for a presentation slide.

CANVAS AND BACKGROUND

- Exactly 3840 x 2160 pixels (16:9).
- Fully transparent background, true alpha channel.
- No background colour, no plate, no frame, no border.
- Every drawn element is a single ink colour: very dark charcoal grey #1A1A1A.
- No other colours anywhere.

MAP

- Equirectangular projection, longitude -180 at the left edge to +180 at the right, latitude +84 at the top to -60 at the bottom.
- Map body spans x = 154 px to x = 3686 px, vertically centred, height 1413 px.
- Continents and major islands as thin uniform outlines only, stroke 4 px, same charcoal ink. Coastlines only: no land fill, no country borders, no country names, no gridlines, no equator or meridian lines, no rivers, no compass, no scale bar, no legend, no icons, no flags.

TITLE - render this exact line and nothing else in the top band:

Around the World with Numeratis

Centred horizontally, at y = 200 px, bold condensed sans-serif, cap height 200 px, same charcoal ink. Top band y = 0 to 400 px is reserved for it: no other element



Around the World with Numeratis



Skill: "IIm-fusion"

IIm-ensemble high: read <https://docs.numer.ai/> and then explain Numerai's stake-weighted meta-model to a 5-yo

20:59 ✓✓

deepseek-v4-flash:

Numerai's stake-weighted meta-model is like a game where many people guess stock prices, and each person bets special coins to show they're confident. The model makes one big guess by paying more attention to people who bet more coins, because they might be extra smart. This creates a much better guess than any single person could make on their own.

mimo-v2.5:

Imagine you and your friends are guessing where a ball is hidden. You each put your own allowance money on your guess — the more money you bet, the more you think you're right. Numerai does the same thing: thousands of data scientists submit predictions about the stock market, and Numerai weighs the people who bet more of their own money (called "staking") more heavily. Then, all those predictions get blended together into one big super-guess (the "meta-model"). This super-guess is usually way smarter than any one person's guess alone — kind of like how the whole class guessing jellybeans in a jar gets closer than any single kid.

minimax-m3:

Final answer (judge):

Imagine you and all your friends are in a classroom, and there's a huge jar of jellybeans on the teacher's desk. Everyone wants to guess how many jellybeans are inside. But here's the fun part: each of you has some shiny gold stars that you really, really like.

>

Before the teacher looks, you can **bet some of your stars** on your guess to show how sure you are. If your guess is good, you win even more stars — but if it's bad, you lose the stars you bet. So you only bet a lot of stars when you're *super* confident!

>

Now, the teacher doesn't just pick one kid's guess. Instead, she takes **everyone's guesses** and mixes them together into one "**super guess**." But she doesn't mix them evenly — she listens much more to the guesses that have lots of stars bet on them. Why? Because those friends were so sure they were willing to risk their own stars. We call that having "**skin in the game**."

>

The way the teacher blends all the guesses, paying extra attention to the ones with bigger star bets, is like a computer building a **meta-model**. And guess what? The teacher's super guess is almost always better than any single kid's guess, because it uses all the brain-power in the room!

>

This is exactly how Numerai works. Grown-ups don't guess jellybeans; they guess if company shares will go up or down. Instead of gold stars, they bet special internet coins (called "crypto"). A computer collects all those guesses, gives more weight to the guesses that have more coins bet on them — that's the "stake" — and blends them into one super-smart prediction. And that super prediction helps make really good money decisions.

>

🌟 That's the **stake-weighted meta-model**!

Skill: “agent-pipeline”

Workhorse Models
(e.g. GPT 6 Luna)

Hermes Alpha Manager

Understood — one more reviewer gate after rollout, and no human gate until the final HTMLs land. Here's the complete workflow:

The full pipeline (no human stops until the end)

1. Design memo — snr-dev

Produces the full design spec (duration 8h/9h, Leaflet grayscale map, lat/lon fields, closing-time checks, test plan, rollout plan)

2. Design review — reviewer

Checks the design for soundness, completeness, consistency with your rulings. Fix loop with snr-dev if needed, until PASS

3. Implement — jnr-dev

Code changes exactly per approved design + tests + installed-skill sync

4. Implementation review — reviewer

Checks code against design. Fix loop with jnr-dev until PASS

5. Verify — jnr-dev

9-point end-state check (tests, no-network rules except map tiles, diff scope, sync)

6. Data update — jnr-research

Geocodes venues (lat/lon), checks closing times / 5pm–6pm allowance for NYC + Montreal shortlists, sources every fact

7. Data review — reviewer

Are the new facts sourced, complete, correct? Fix loop with jnr-research until PASS

8. Rollout — jnr-research

Re-scores with 8h/9h costs, republishes bundles, regenerates the 3 summaries with the grayscale map, sanity-checks output

9. Rollout review — reviewer

Final gate on the actual published bundles + summaries: costs correct at 8h, map present and functional, provenance intact, tiers sane. Fix loop until PASS

10. Deliver — final HTMLs to you ← the only human gate

Premium
(e.g. GPT 6 Sol)

Mid-tier
(e.g. GLM
5.3 Flash)

Skill: “agent-pipeline”

Next stage (WIP):

Using “System One”
model (e.g. Jev, This-
That) to assign
models and skills

Pinned

Flock.io 
@flock_io

Introducing THIS / THAT Model 1.0. Trained by Flock.

Not another chatbot. A best-in-class open-source model built to decide.

94.1% success rate on our 68-question decision test, outperforming GPT-5.6, Deepseek-4.1-flash, Kimi K3, Jev, and more.

Just 1.9B parameters. Runs locally. Measured decision latency of ~31 ms on an RTX 5080 laptop GPU.

We're putting decentralized AI to work: giving developers and AI agents decision-making capabilities they can run on their own hardware.

Your model. Your machine. Not every decision needs the cloud.

THIS / THAT
MODEL
trained by **Flock**

